

# Homework 7. Naive Bayesian - Spam or Ham

*Double Click here to edit this cell*

- Name: 전기범
- Student ID: 201703091
- Submission date: 2019/06/03

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
import re
from sklearn.metrics import confusion_matrix, classification_report
from collections import Counter
```

```
In [2]: messages = pd.read_csv("spam_utf8.csv")
```

**We have 5572 text messages (747 spams or 4825 hams)**

## Problem 1 (5 pts): Most Common Words

- Use the following for word-splitting:

```
re.findall("[a-z0-9' ]+", text.lower())
```

- Find 20 most common spam words

```
In [3]: messages.head()
```

Out[3]:

	category	text
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
In [4]: spam_text=np.sum(messages.loc[messages['category']=='spam'].text.values+ '\n')
# print(spam_text)
spam=re.findall("[a-z0-9' ]+", spam_text.lower())
# print(spam)
spam_cnt=Counter(spam)

print([(key,value) for key,value in sorted(dict(spam_cnt).items(),
key =lambda x:x[1], reverse=True)[:20])
# YOUR CODE MUST BE HERE

[('to', 688), ('a', 378), ('call', 355), ('you', 290), ('your', 264), ('free', 224), ('2', 206), ('the', 206), ('for', 204), ('now', 202), ('or', 188), ('u', 169), ('txt', 163), ('is', 158), ('on', 144), ('ur', 144), ('4', 137), ('have', 135), ('from', 131), ('mobile', 127)]
```

**Your output should be like the following:**

```
[('to', 688), ('a', 378), ('call', 355), ('you', 290), ('your', 264), ('free', 224), ('2', 206), ('the', 206), ('for', 204), ('now', 202), ('or', 188), ('u', 169), ('txt', 163), ('is', 158), ('on', 144), ('ur', 144), ('4', 137), ('have', 135), ('from', 131), ('mobile', 127)]
```

- Find 20 most common ham words

```
In [5]: ham_text="\n".join(str(h) for h in messages.loc[messages['category']=='ham'].text.values)
ham=re.findall("[a-z0-9' ]+", ham_text.lower())
ham_cnt=Counter(ham)
print([(key,value) for key,value in sorted(dict(ham_cnt).items(), key =lambda x:x[1], reverse=True)[:20])
# YOUR CODE MUST BE HERE

[('i', 2281), ('you', 1836), ('to', 1518), ('the', 1100), ('a', 1046), ('u', 980), ('and', 843), ('in', 802), ('me', 758), ('my', 734), ('is', 709), ('it', 602), ('of', 513), ('for', 490), ('that', 486), ('have', 430), ('but', 427), ('so', 422), ('are', 408), ('your', 407)]
```

**Your output should be like the following:**

```
[('i', 2281), ('you', 1836), ('to', 1518), ('the', 1100), ('a', 1046), ('u', 980), ('and', 843), ('in', 802), ('me', 758), ('my', 734), ('is', 709), ('it', 602), ('of', 513), ('for', 490), ('that', 486), ('have', 430), ('but', 427), ('so', 422), ('are', 408), ('your', 407)]
```

## Problem 2 (10 pts): My First Simplest Spam Filter

I designed a very simple spam classifier:

- If the text contains one of the spam words, it is a spam
- otherwise, it is a ham

Note that:

- Use the following for word-splitting:

```
re.findall("[a-z0-9' ]+", text.lower())
```

For each spam words:

```
['call']  
['call', 'free']  
['call', 'free', 'txt']
```

- Compute confusion matrix, precision, recall, f1-score using all data as a test dataset

```
In [6]: import re  
texts = [str(text) for text in messages['text']]  
  
class MySimplestSpamFilter(object):  
  
    def fit(self, spam_words):  
        # FILL OUT  
        self.spam_words = spam_words  
  
    def predict(self, texts):  
        res = []  
        # FILL OUT  
        for text in texts:  
            words = [word for word in re.findall("[a-z0-9' ]+", text  
.lower())]  
            if set(self.spam_words) & set(words):  
                res.append('spam')  
            else:  
                res.append('ham')  
        return res
```

```
In [7]: # RUN THIS CELL
model = MySimplestSpamFilter()
model.fit(['call'])

y_pred = model.predict(texts)
y_true = messages['category']

print(confusion_matrix(y_true, y_pred, labels=['spam', 'ham']).T)
print(classification_report(y_true, y_pred, labels=['spam', 'ham']))
```

```
[[ 328  218]
 [ 419 4607]]

              precision    recall  f1-score   support

      spam         0.60         0.44         0.51         747
      ham         0.92         0.95         0.94        4825

   micro avg         0.89         0.89         0.89        5572
   macro avg         0.76         0.70         0.72        5572
weighted avg         0.87         0.89         0.88        5572
```

## YOUR OUTPUT SHOULD BE:

```
[[ 328  218]
 [ 419 4607]]

              precision    recall  f1-score   support

      spam         0.60         0.44         0.51         747
      ham         0.92         0.95         0.94        4825

avg / total         0.87         0.89         0.88        5572
```

```
In [8]: # RUN THIS CELL
model = MySimplestSpamFilter()
model.fit(['call', 'free'])

y_pred = model.predict(texts)
y_true = messages['category']

print(confusion_matrix(y_true, y_pred, labels=['spam', 'ham']).T)
print(classification_report(y_true, y_pred, labels=['spam', 'ham'])
)
```

```
[[ 438  265]
 [ 309 4560]]

              precision    recall  f1-score   support

      spam         0.62         0.59         0.60         747
      ham         0.94         0.95         0.94        4825

   micro avg         0.90         0.90         0.90        5572
   macro avg         0.78         0.77         0.77        5572
weighted avg         0.89         0.90         0.90        5572
```

## YOUR OUTPUT SHOULD BE:

```
[[ 438  265]
 [ 309 4560]]

              precision    recall  f1-score   support

      spam         0.62         0.59         0.60         747
      ham         0.94         0.95         0.94        4825

avg / total         0.89         0.90         0.90        5572
```

```
In [9]: # RUN THIS CELL
model = MySimplestSpamFilter()
model.fit(['call', 'free', 'txt'])

y_pred = model.predict(texts)
y_true = messages['category']

print(confusion_matrix(y_true, y_pred, labels=['spam', 'ham']).T)
print(classification_report(y_true, y_pred, labels=['spam', 'ham'])
)
```

```
[[ 532  277]
 [ 215 4548]]

              precision    recall  f1-score   support

    spam           0.66       0.71       0.68         747
    ham           0.95       0.94       0.95        4825

   micro avg       0.91       0.91       0.91       5572
   macro avg       0.81       0.83       0.82       5572
weighted avg       0.92       0.91       0.91       5572
```

## YOUR OUTPUT SHOULD BE:

```
[[ 532  277]
 [ 215 4548]]

              precision    recall  f1-score   support

    spam           0.66       0.71       0.68         747
    ham           0.95       0.94       0.95        4825

avg / total       0.92       0.91       0.91       5572
```

## Problem 3 (5 pts): My Another Simplest Spam Filter

I designed another very simple spam classifier:

- If the text contains **all of the spam words**, it is a spam
- otherwise, it is a ham

Note that:

- Use the following for word-splitting:

```
re.findall("[a-z0-9' ]+", text.lower())
```

For this set of spam words:

```
['call', 'free', 'txt']
```

- Compute confusion marix, precision, recall, f1-score using all data as a test dataset

```
In [10]: import re

class MyAnotherSimplestSpamFilter(object):
    def fit(self, spam_words):
        self.spam_words = spam_words

    def predict(self, texts):
        res = []
        for text in texts:
            words = [word for word in re.findall("[a-z0-9' ]+", text
.lower())]
            if set(self.spam_words) <= set(words):
                res.append('spam')
            else:
                res.append('ham')
        return res
```

```
In [11]: # RUN THIS CELL
model = MyAnotherSimplestSpamFilter()
model.fit(['call', 'free', 'txt'])

y_pred = model.predict(texts)
y_true = messages['category']

print(confusion_matrix(y_true, y_pred, labels=['spam', 'ham']).T)
print(classification_report(y_true, y_pred, labels=['spam', 'ham']))
```

```
[[ 6  0]
 [741 4825]]
```

	precision	recall	f1-score	support
spam	1.00	0.01	0.02	747
ham	0.87	1.00	0.93	4825
micro avg	0.87	0.87	0.87	5572
macro avg	0.93	0.50	0.47	5572
weighted avg	0.88	0.87	0.81	5572

#### YOUR OUTPUT SHOULD BE:

```
[[ 6  0]
 [741 4825]]
```

	precision	recall	f1-score	support
spam	1.00	0.01	0.02	747
ham	0.87	1.00	0.93	4825
avg / total	0.88	0.87	0.81	5572

## Problem 4. Discussion (5 pts)

- From the Problem 2, 3 experiments, discuss about precision and recall trade-off.
- WRITE HERE (To edit, double click this cell)

recall은 내가 모델로부터 얻은 결과중에 얼마나 정답이 포함되었는가 여부이고, precision은 모델에서 얻은 결과중에 얼마나 정답일 가능성이 높은가이다. 두 가지 용어는 서로 trade off 관계이다.

단어의 수가 증가함에 따라 Spam 판정의 확률이 높아지는것을 보아 Spam Filter를 위해 단어의 수를 증가시키면 더 정확도가 높은 필터링이 되는것을 확인 할 수 있었습니다



## Problem 5. SMS Spam Filter (20 pts)

- Using sklearn modules, implement your SMS spam filter
- Print confusion matrix, classification report, **f0.5 score**
- This homework will be graded based on **f0.5 score**

```
In [12]: # DO NO EDIT THIS CELL
import numpy as np

np.random.seed(0)
```

```
In [13]: # YOU CODE MUST BE HERE

# import whatever module as you need
from sklearn.naive_bayes import BernoulliNB
from sklearn.metrics import fbeta_score
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer

# You must have 30% stratified sample from dataset for test dataset
like the following:

vectorizer = CountVectorizer(stop_words = "english", lowercase=True
, binary=False) # default actually
vecs = vectorizer.fit_transform(texts) # it is a sparse matrix rep
resentation

targets = messages["category"].apply(lambda c: 1 if c == 'spam' els
e 0)

X_train, X_test, y_train, y_test = train_test_split(vecs, targets,
stratify = targets, test_size = 0.3)

bernoulli_NB = BernoulliNB()
bernoulli_NB.fit(X_train, y_train)

y_pred = bernoulli_NB.predict(X_test) # default binarize=0.0

print(confusion_matrix(y_test, y_pred, labels=[1, 0]).T)
print(classification_report(y_test, y_pred, labels=[1, 0]))
beta = 0.5
print('f_{0.5} score is {0:6.2f}%'.format(beta, fbeta_score(y_test, y_p
red, beta = beta)*100))
```

```
[[ 183    8]
 [  41 1440]]
```

	precision	recall	f1-score	support
1	0.96	0.82	0.88	224
0	0.97	0.99	0.98	1448
micro avg	0.97	0.97	0.97	1672
macro avg	0.97	0.91	0.93	1672
weighted avg	0.97	0.97	0.97	1672

f<sub>0.5</sub> score is 92.61%

## Problem 6. MyBernoulliNaiveBayesClassifier (40 pts)

- Implement your own Bernoulli Naive Bayes Classifier

```
In [ ]: class MyBernoulliNaiveBayesClassifier(object):
        def fit(self, X, Y, smoothing=1):
            # FILL OUT

        def predict(self, X):
            res = []
            # FILL OUT
            return res
```

```
In [ ]: # RUN THIS CELL

import numpy as np
np.random.seed(0)
X = np.random.randint(2, size=(6, 100))
y = np.array([1, 2, 3, 3, 1, 0])
my_nbc = MyBernoulliNaiveBayesClassifier()
my_nbc.fit(X, y)

y_pred = my_nbc.predict(X)

print(confusion_matrix(y, y_pred).T)
print(classification_report(y, y_pred))
```

**YOUR OUTPUT MUST BE LIKE THE FOLLOWING:**

```
[[1 0 0 0]
 [0 2 0 0]
 [0 0 1 0]
 [0 0 0 2]]

              precision    recall  f1-score   support

    0           1.00         1.00         1.00           1
    1           1.00         1.00         1.00           2
    2           1.00         1.00         1.00           1
    3           1.00         1.00         1.00           2

 avg / total           1.00         1.00         1.00           6
```

## Ethics:

If you cheat, you will get negative of the total points. If the homework total is 22 and you cheat, you get -22.

## What to submit

- Run all cells
- Goto "File -> Print Preview"
- Print the page
- Submit in class
- No late homeworks accepted
- Your homework will be graded on the basis of correctness and programming skills

**Deadline: 6/3**